

```
In [1]: 1 !kaggle datasets download jacopoferretti/bbc-articles-dataset
        2 !unzip bbc-articles-dataset.zip
```

Dataset URL: <https://www.kaggle.com/datasets/jacopoferretti/bbc-articles-dataset> (<https://www.kaggle.com/datasets/jacopoferretti/bbc-articles-dataset>)

License(s): CC0-1.0

Downloading bbc-articles-dataset.zip to /content

0% 0.00/1.81M [00:00<?, ?B/s]

100% 1.81M/1.81M [00:00<00:00, 130MB/s]

Archive: bbc-articles-dataset.zip

inflating: bbc_text_cls.csv

```
In [2]: 1 import pandas as pd
```

```
In [3]: 1 df = pd.read_csv("/content/bbc_text_cls.csv")
```

```
In [4]: 1 df.head()
```

Out [4]:

	text	labels
0	Ad sales boost Time Warner profit\n\nQuarterly...	business
1	Dollar gains on Greenspan speech\n\nThe dollar...	business
2	Yukos unit buyer faces loan claim\n\nThe owner...	business
3	High fuel prices hit BA's profits\n\nBritish A...	business
4	Pernod takeover talk lifts Domecq\n\nShares in...	business

Pre-Processing Text

```
In [5]: 1 import re
        2 import nltk
        3 from nltk.corpus import stopwords
        4 from nltk.tokenize import word_tokenize
        5 from nltk.stem import WordNetLemmatizer
        6 from bs4 import BeautifulSoup
        7 nltk.download('punkt')
        8 nltk.download('stopwords')
        9 nltk.download('wordnet')
```

[nltk_data] Downloading package punkt to /root/nltk_data...

[nltk_data] Unzipping tokenizers/punkt.zip.

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

[nltk_data] Downloading package wordnet to /root/nltk_data...

Out [5]: True

```

In [6]: 1 def preprocess_text(text):
2         # 1. Lowercasing
3         text = text.lower()
4
5         # 2. Remove HTML Tags
6         text = BeautifulSoup(text, 'html.parser').get_text()
7
8         # 3. Remove URLs
9         text = re.sub(r'https?://\S+|www\.\S+', '', text)
10
11        # 4. Remove emojis
12        emoji_pattern = re.compile("[
13                                     "\U0001F600-\U0001F64F" # Emoti
14                                     "\U0001F300-\U0001F5FF" # Misce
15                                     "\U0001F680-\U0001F6FF" # Trans
16                                     "\U0001F700-\U0001F77F" # Alche
17                                     "\U0001F780-\U0001F7FF" # Geome
18                                     "\U0001F800-\U0001F8FF" # Suppl
19                                     "\U0001F900-\U0001F9FF" # Suppl
20                                     "\U0001FA00-\U0001FA6F" # Chess
21                                     "\U0001FA70-\U0001FAFF" # Symbo
22                                     "\U00002702-\U000027B0" # Dingb
23                                     "\U000024C2-\U0001F251"
24                                     "]" +, flags=re.UNICODE)
25        text = emoji_pattern.sub('', text)
26
27        # 5. Remove punctuation and numbers
28        text = re.sub(r'^a-z\s]', '', text)
29
30        # 6. Tokenization
31        tokens = word_tokenize(text)
32
33        # 7. Remove stopwords and single character tokens
34        stop_words = set(stopwords.words('english'))
35        tokens = [token for token in tokens if token not in stop_wo
36
37        # 8. Lemmatization
38        lemmatizer = WordNetLemmatizer()
39        tokens = [lemmatizer.lemmatize(token) for token in tokens]
40
41        # Joining tokens back into a sentence
42        preprocessed_text = ' '.join(tokens)
43
44        return preprocessed_text

```

```

In [7]: 1 df['text'] = df['text'].apply(preprocess_text)

```

In [8]:

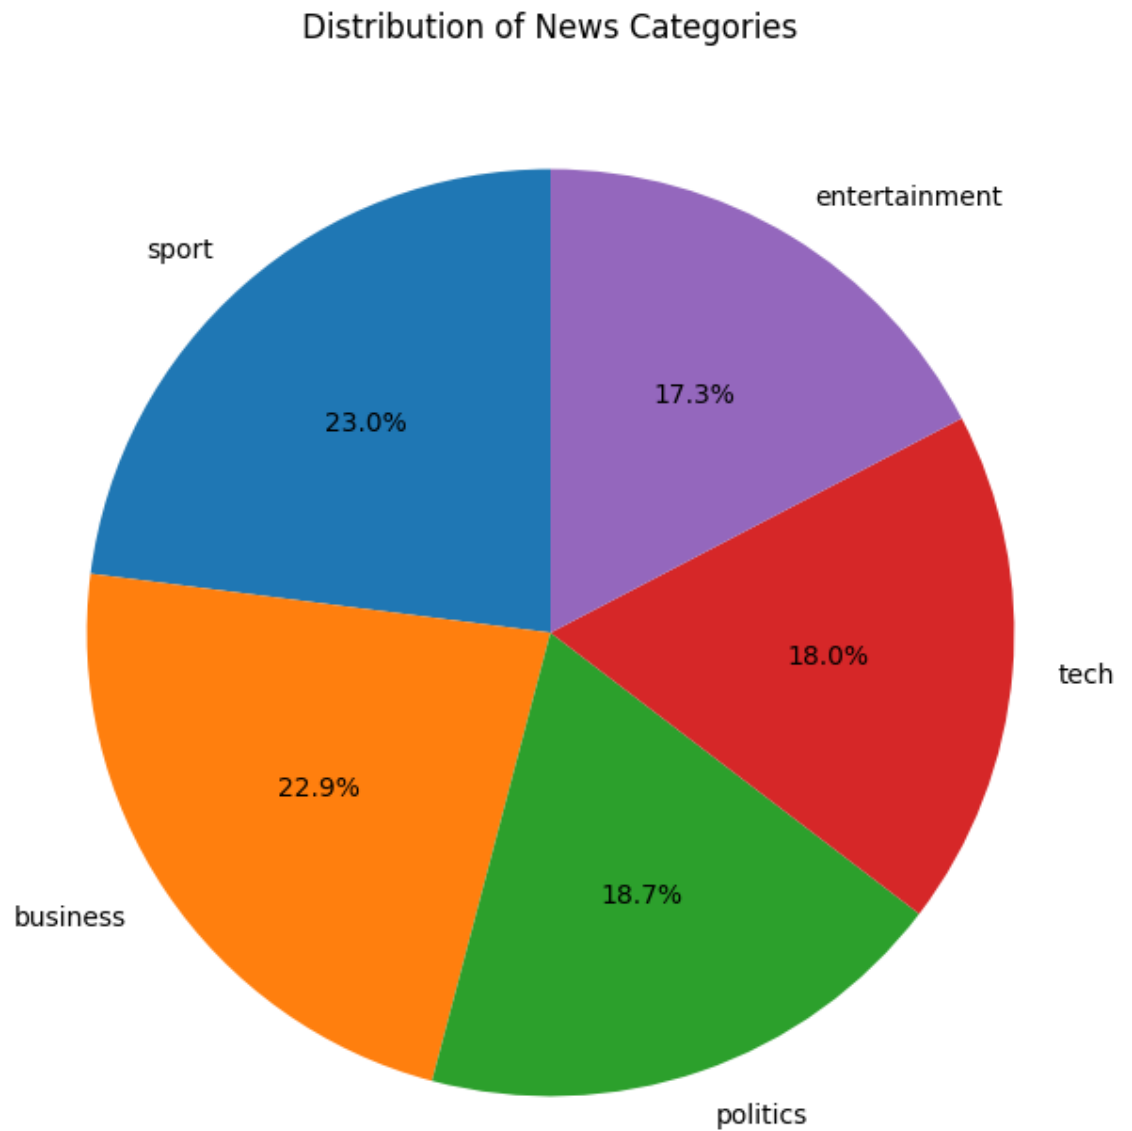
1	df
---	----

Out [8]:


	text	labels
0	ad sale boost time warner profit quarterly pro...	business
1	dollar gain greenspan speech dollar hit highes...	business
2	yukos unit buyer face loan claim owner embattl...	business
3	high fuel price hit ba profit british airway b...	business
4	pernod takeover talk lift domecq share uk drin...	business
...
2220	bt program beat dialler scam bt introducing tw...	tech
2221	spam email tempt net shopper computer user acr...	tech
2222	careful code new european directive could put ...	tech
2223	u cyber security chief resigns man making sure...	tech
2224	losing online gaming online role playing game ...	tech

2225 rows × 2 columns

```
In [9]: 1 # @title Distribution of News Categories
2
3 import matplotlib.pyplot as plt
4
5 category_counts = df['labels'].value_counts()
6
7 plt.figure(figsize=(8, 8))
8 plt.pie(category_counts, labels=category_counts.index, autopct=
9 _ = plt.title('Distribution of News Categories')
```



```
In [10]: 1 # @title labels
2
3 from matplotlib import pyplot as plt
4 import seaborn as sns
5 df.groupby('labels').size().plot(kind='barh', color=sns.palette
6 plt.gca().spines[['top', 'right',]].set_visible(False)
```



Classification Using Machine learning model

Convert Text to Numerical Format (Tokenization)

1. Bag-of-Words (BoW)

- Counts the occurrence of each word in a document.
- Generates sparse, high-dimensional vectors where each word is a feature.
- Pros: Simple and effective for many applications.
- Cons: Ignores word order and meaning.
- **Implementation:** CountVectorizer in Scikit-Learn.

2. Term Frequency-Inverse Document Frequency (TF-IDF)

- Weighs word importance based on how frequently it appears in a document vs. across all documents.
- Useful for emphasizing significant words that appear frequently in a document but

not commonly across documents.

- Pros: Reduces the impact of common but less informative words.
 - Cons: Still ignores word order and context.
 - **Implementation:** `TfidfVectorizer` in Scikit-Learn.
-

3. n-grams

- Extracts contiguous sequences of `n` words to capture basic word order and simple context.
 - Common choices: bigrams (`n=2`), trigrams (`n=3`).
 - Pros: Adds some contextual information beyond individual words.
 - Cons: Increases feature space and can add noise if `n` is too large.
 - **Implementation:** `CountVectorizer` or `TfidfVectorizer` with `ngram_range` parameter in Scikit-Learn.
-

4. Word Embeddings

- Maps words to dense, low-dimensional vectors that capture semantic meaning.
 - Popular pre-trained embeddings:
 - **Word2Vec:** Trained on word co-occurrences; represents semantic relationships.
 - **GloVe:** Trained on global word-word co-occurrence statistics.
 - Pros: Captures meaning and relationships between words.
 - Cons: Requires pre-trained embeddings and doesn't inherently capture sentence structure.
 - **Implementation:** Using libraries like `gensim` or loading pre-trained embeddings in `PyTorch/Keras`.
-

5. Pre-trained Transformer Embeddings

- Use pre-trained models like BERT, DistilBERT, or RoBERTa to extract embeddings for text.
- Contextualized embeddings capture word meaning based on context.
- Pros: State-of-the-art embeddings for most NLP tasks; captures contextual meaning.
- Cons: Requires more computation and memory.
- **Implementation:** Hugging Face Transformers library.

```
In [41]: 1 # @title Splitting Data
          2
          3 from sklearn.model_selection import train_test_split
          4
          5 X_train, X_test, y_train, y_test = train_test_split(df['text'],
```

```
In [53]: 1 from sklearn.preprocessing import LabelEncoder
2
3 label_encoder = LabelEncoder()
4 y_train_encoded = label_encoder.fit_transform(y_train)
5 y_test_encoded = label_encoder.transform(y_test)
```

```
In [43]: 1 # @title 1. Bag-of-Words (BoW)
2
3 from sklearn.feature_extraction.text import CountVectorizer
4
5 bow_vectorizer = CountVectorizer()
6
7 X_train_bow = bow_vectorizer.fit_transform(X_train)
8 X_test_bow = bow_vectorizer.transform(X_test)
```

```
In [58]: 1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.metrics import accuracy_score, classification_report
3
4 nb_bow = MultinomialNB()
5 nb_bow.fit(X_train_bow, y_train)
6
7 # Predict on the test data
8 y_pred_bow = nb_bow.predict(X_test_bow)
9
10 # Evaluate the BoW model
11 print("Bag-of-Words Model Evaluation")
12 print("Accuracy:", accuracy_score(y_test, y_pred_bow))
13 print("Classification Report:\n", classification_report(y_test,
```

Bag-of-Words Model Evaluation

Accuracy: 0.9685393258426966

Classification Report:

	precision	recall	f1-score	support
business	0.97	0.96	0.96	115
entertainment	0.99	0.93	0.96	72
politics	0.93	0.97	0.95	76
sport	1.00	0.99	1.00	102
tech	0.95	0.99	0.97	80
accuracy			0.97	445
macro avg	0.97	0.97	0.97	445
weighted avg	0.97	0.97	0.97	445

```
In [49]: 1 # @title 2. Term Frequency-Inverse Document Frequency (TF-IDF)
2 from sklearn.feature_extraction.text import TfidfVectorizer
3 from sklearn.metrics import accuracy_score, classification_report
4
5 tfidf_vectorizer = TfidfVectorizer()
6
7 # Transform the text data into TF-IDF features
8 X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
9 X_test_tfidf = tfidf_vectorizer.transform(X_test)
10
11 # Train a Naive Bayes classifier on TF-IDF features
12 nb_tfidf = MultinomialNB()
13 nb_tfidf.fit(X_train_tfidf, y_train)
14
15 # Predict on the test data
16 y_pred_tfidf = nb_tfidf.predict(X_test_tfidf)
17
18 # Evaluate the TF-IDF model
19 print("TF-IDF Model Evaluation")
20 print("Accuracy:", accuracy_score(y_test, y_pred_tfidf))
21 print("Classification Report:\n", classification_report(y_test,
```

TF-IDF Model Evaluation

Accuracy: 0.9617977528089887

Classification Report:

	precision	recall	f1-score	support
business	0.97	0.95	0.96	115
entertainment	0.98	0.90	0.94	72
politics	0.90	0.97	0.94	76
sport	0.99	0.99	0.99	102
tech	0.95	0.99	0.97	80
accuracy			0.96	445
macro avg	0.96	0.96	0.96	445
weighted avg	0.96	0.96	0.96	445

```
In [50]: 1 text = '''
2 wru proposes season overhaul welsh rugby union want restructure
3 start celtic league october followed heineken cup february march
4 twomonth period away home international match wru chairman david
5 club country added feel sure spectator interest would respond im
6 timetable currently operation equally suspect sponsor would pref
7 would also enjoy increased exposure moving six nation traditiona
8 greater interest game generally provide increased skill competit
9 international rugby board next month four plan drawn independent
10 early day number caveat associated least revenue broadcaster ext
11 '''
```

```
In [67]: 1 # @title Tfidf Prediction
2
3 nb_tfidf.predict(tfidf_vectorizer.transform([text]))[0]
```

Out [67]: 'sport'


```
In [68]: 1 # @title BOW Prediction
          2
          3 nb_bow.predict(bow_vectorizer.transform([text]))[0]
```

```
Out[68]: 'sport'
```